

## **COMPLETE GUIDE: Building QEMU 4.2.1 with vGPU Stub Device for XCP-ng**

Author: Bren Max

Date: January 8, 2025

Tested On: XCP-ng 8.2 (Xen 4.13, QEMU 4.2.1)

Status: VERIFIED WORKING ✓

This guide provides the EXACT steps used to successfully build QEMU with a custom vGPU stub device that appears as a PCI device in guest VMs.

**IMPORTANT:** These steps have been tested and work. Follow them exactly.

### **SECTION 1: WHAT YOU'LL BUILD**

A custom PCI device that:

- Shows up in guest VMs as: "Processing accelerator: Red Hat, Inc. Device 1111"
- Can be detected with lspci
- Has a 4KB MMIO region for future communication
- Is a foundation for more complex vGPU implementations

Device Details:

Vendor ID: 0x1AF4 (Red Hat, Inc.)

Device ID: 0x1111 (Custom)

Class: Processing Accelerator (Co-processor)

BAR0: 4KB memory-mapped I/O region

### **SECTION 2: PREREQUISITES**

System Requirements:

- XCP-ng 8.2 or compatible Xen system
- At least 2GB free disk space
- Internet connection for downloads
- Root access

Required Packages:

```
yum install -y gcc gcc-c++ make python2 glib2-devel pixman-devel  
zlib-devel bzip2 wget git centos-release-scl
```

Enable Developer Toolset (for modern GCC):

```
yum install -y devtoolset-11  
scl enable devtoolset-11 bash
```

Verify GCC version:

```
gcc --version  
# Should show GCC 11.x or higher
```

**CRITICAL:** You must enable devtoolset-11 before building, or you'll get errors!

### **SECTION 3: DOWNLOAD AND EXTRACT QEMU**

#### **Step 1: Create Clean Working Directory**

We start fresh to avoid issues with previous build attempts.

```
mkdir -p ~/qemu-clean-build  
cd ~/qemu-clean-build
```

Step 2: Download QEMU 4.2.1 (exact version XCP-ng uses)

wget <https://download.qemu.org/qemu-4.2.1.tar.xz>

Verify download:

ls -lh qemu-4.2.1.tar.xz

```
10:18 xcp-ng-syovfxoz qemu-clean-build # ls -lh qemu-4.2.1.tar.xz
-rw-r--r-- 1 root root 60M Oct 26 2024 qemu-4.2.1.tar.xz
10:18 xcp-ng-syovfxoz qemu-clean-build #
```



# Should show ~60MB file

Step 3: Extract

```
tar xf qemu-4.2.1.tar.xz
cd qemu-4.2.1
```

Current directory should be:

/root/qemu-clean-build/qemu-4.2.1

#### SECTION 4: CREATE THE VGPU-STUB DEVICE CODE

Step 1: Create the Device Source File

We create a QEMU 4.2.1-compatible device using the older QOM API.

IMPORTANT: The code below is for QEMU 4.2.1. Do NOT use QEMU 7.x code!

```
cat > hw/misc/vgpu-stub.c << 'VGPU_CODE'
/*
 * VGPU Stub Device for XCP-ng
 *
 * Simple PCI device that appears as a GPU in guest VMs
 * Compatible with QEMU 4.2.1 API
 */

#include "qemu/osdep.h"
#include "hw/pci/pci.h"
#include "hw/hw.h"
#include "hw/pci/msi.h"
#include "qemu/timer.h"
#include "qom/object.h"
#include "qemu/module.h"
#include "sysemu/kvm.h"

#define TYPE_VGPU_STUB "vgpu-stub"
#define VGPU_STUB(obj) OBJECT_CHECK(VGPUStubState, (obj), TYPE_VGPU_STUB)
```

```

typedef struct VGPUStubState {
    PCIDevice parent_obj;

    /* Memory-mapped I/O region */
    MemoryRegion mmio;

    /* Device state */
    uint32_t command_reg;
    uint32_t status_reg;
} VGPUStubState;

/* MMIO read handler */
static uint64_t vgpu_mmio_read(void *opaque, hwaddr addr, unsigned size)
{
    VGPUStubState *s = opaque;
    uint64_t val = 0;

    switch (addr) {
    case 0x00: /* Command register */
        val = s->command_reg;
        break;
    case 0x04: /* Status register */
        val = s->status_reg;
        break;
    default:
        val = 0;
        break;
    }

    return val;
}

/* MMIO write handler */
static void vgpu_mmio_write(void *opaque, hwaddr addr, uint64_t val, unsigned size)
{
    VGPUStubState *s = opaque;

    switch (addr) {
    case 0x00: /* Command register */
        s->command_reg = val;
        /* Set status to indicate command received */
        s->status_reg = 0x1;
        break;
    case 0x04: /* Status register (read-only in real device) */
        /* Ignore writes */
        break;
    default:
        break;
    }
}

static const MemoryRegionOps vgpu_mmio_ops = {
    .read = vgpu_mmio_read,
    .write = vgpu_mmio_write,
    .endianness = DEVICE_LITTLE_ENDIAN,
    .impl = {

```

```

    .min_access_size = 4,
    .max_access_size = 4,
},
};

/* Device realization */
static void vgpu_realize(PCIDevice *pci_dev, Error **errp)
{
    VGPUStubState *s = VGPU_STUB(pci_dev);

    /* Set up PCI config space */
    pci_dev->config[PCI_INTERRUPT_PIN] = 1; /* Interrupt pin A */

    /* Initialize MMIO region (4KB) */
    memory_region_init_io(&s->mmio, OBJECT(s), &vgpu_mmio_ops, s,
        "vgpu-stub-mmio", 4096);
    pci_register_bar(pci_dev, 0, PCI_BASE_ADDRESS_SPACE_MEMORY, &s->mmio);

    /* Initialize device state */
    s->command_reg = 0;
    s->status_reg = 0;
}

/* Device exit */
static void vgpu_exit(PCIDevice *pci_dev)
{
    /* Cleanup if needed */
}

static void vgpu_class_init(ObjectClass *klass, void *data)
{
    DeviceClass *dc = DEVICE_CLASS(klass);
    PCIDeviceClass *k = PCI_DEVICE_CLASS(klass);

    k->realize = vgpu_realize;
    k->exit = vgpu_exit;
    k->vendor_id = 0x1AF4; /* Red Hat */
    k->device_id = 0x1111; /* Custom device ID */
    k->revision = 0x01;
    k->class_id = PCI_CLASS_PROCESSOR_CO; /* Co-processor */

    set_bit(DEVICE_CATEGORY_MISC, dc->categories);
    dc->desc = "Virtual GPU Stub Device";
}

static const TypeInfo vgpu_info = {
    .name      = TYPE_VGPU_STUB,
    .parent    = TYPE_PCI_DEVICE,
    .instance_size = sizeof(VGPUStubState),
    .class_init = vgpu_class_init,
    .interfaces = (InterfaceInfo[]) {
        { INTERFACE_CONVENTIONAL_PCI_DEVICE },
        { },
    },
};

static void vgpu_register_types(void)

```

```
{
    type_register_static(&vgpu_info);
}
```

```
type_init(vgpu_register_types)
VGPU_CODE
```

```

GNU nano 2.3.3
root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
File: hw/misc/vgpu-stub.c

/*
 * VGPU Stub Device for XCP-ng
 *
 * Simple PCI device that appears as a GPU in guest VMs
 */
#include "qemu/osdep.h"
#include "hw/pci/pci.h"
#include "hw/hw.h"
#include "hw/pci/msl.h"
#include "qemu/timer.h"
#include "qom/object.h"
#include "qemu/module.h"
#include "sysemu/kvm.h"

#define TYPE_VGPU_STUB "vgpu-stub"
#define VGPU_STUB(obj) OBJECT_CHECK(VGPUStubState, (obj), TYPE_VGPU_STUB)

typedef struct VGPUStubState {
    PCIDevice parent_obj;

    /* Memory-mapped I/O region */
    MemoryRegion mmio;

    /* Device state */
    uint32_t command_reg;
    uint32_t status_reg;
} VGPUStubState;

/* MMIO read handler */
static uint64_t vgpu_mmio_read(void *opaque, hwaddr addr, unsigned size)
{
    VGPUStubState *s = opaque;
}

```

File: hw/misc/vgpu-stub.c

File Edit View Insert Block Comment Search Help

Get Help WriteOut Read File Cut Text Cur Pos

Exit Justify Where Is Prev Page Uncut Text To Spell

[ Read 136 lines ]

136 hw/misc/vgpu-stub.c

### Step 2: Verify File Created

```
ls -lh hw/misc/vgpu-stub.c
```

```

root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
ls -lh hw/misc/vgpu-stub.c
-rw-r--r-- 1 root root 3.2K Dec 21 08:06 hw/misc/vgpu-stub.c

```

# Should show ~4KB file

```
wc -l hw/misc/vgpu-stub.c
```

```

root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
wc -l hw/misc/vgpu-stub.c
136 hw/misc/vgpu-stub.c

```

# Should show 136 lines

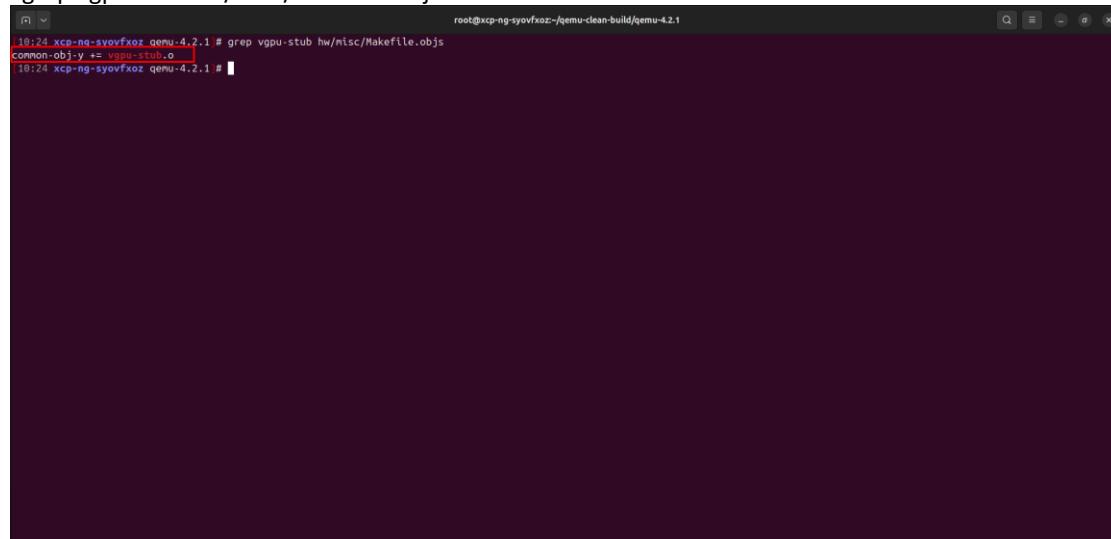
### Step 3: Update Build System (Makefile.objs)

Add vgpu-stub to the build by editing hw/misc/Makefile.objs

```
sed -i '/common-obj-$(CONFIG_EDU) += edu.o/a common-obj-y += vgpu-stub.o' \
hw/misc/Makefile.objs
```

Verify it was added:

```
grep vgpu-stub hw/misc/Makefile.objs
```

A terminal window titled 'root@xcp-ng-syovfxoz:~/.qemu-clean-build/qemu-4.2.1'. It contains the command 'grep vgpu-stub hw/misc/Makefile.objs' and its output: 'common-obj-y += vgpu-stub.o'. The line 'common-obj-y += vgpu-stub.o' is highlighted with a red box.

Expected output:

```
common-obj-y += vgpu-stub.o
```

Check context (should be after edu.o):

```
sed -n '8,12p' hw/misc/Makefile.objs
```

Expected output:

```
common-obj-$(CONFIG_PCI_TESTDEV) += pci-testdev.o
common-obj-$(CONFIG_EDU) += edu.o
common-obj-y += vgpu-stub.o
common-obj-$(CONFIG_PCA9552) += pca9552.o
```

## **SECTION 5: CONFIGURE QEMU**

### CRITICAL Configuration Notes:

- We use --disable-xen because xen-devel headers may not be installed
- We use --disable-werror to avoid warnings blocking the build
- We use --python=/usr/bin/python2 because QEMU 4.2.1 needs Python 2
- We build x86\_64-softmmu which can run both 32-bit and 64-bit VMs

### Run Configure:

```
./configure \
--target-list=x86_64-softmmu \
--prefix=/usr \
--disable-xen \
--enable-kvm \
--disable-werror \
--python=/usr/bin/python2
```

Expected Output (check these lines):

```
Install prefix  /usr
target list    x86_64-softmmu
xen support   no
KVM support   yes
python        /usr/bin/python2 -B (2.7.5)
```

If configure fails, check:

1. GCC version: gcc --version (should be 11.x)
2. Python2 exists: which python2
3. Required libraries: yum list installed | grep -E "glib2-devel|pixman-devel"

## SECTION 6: BUILD QEMU

Step 1: Start Build

---

```
make -j$(nproc) 2>&1 | tee build.log
```

This will take 10-20 minutes depending on your system.

The tee command saves output to build.log for troubleshooting.

Expected Progress:

```
[Start] Generating configuration files
[Middle] Compiling hundreds of .c files
      Look for: CC hw/misc/vgpu-stub.o (around line 1489 in our test)
[End] LINK x86_64-softmmu/qemu-system-x86_64
```

### Step 2: Check Build Status

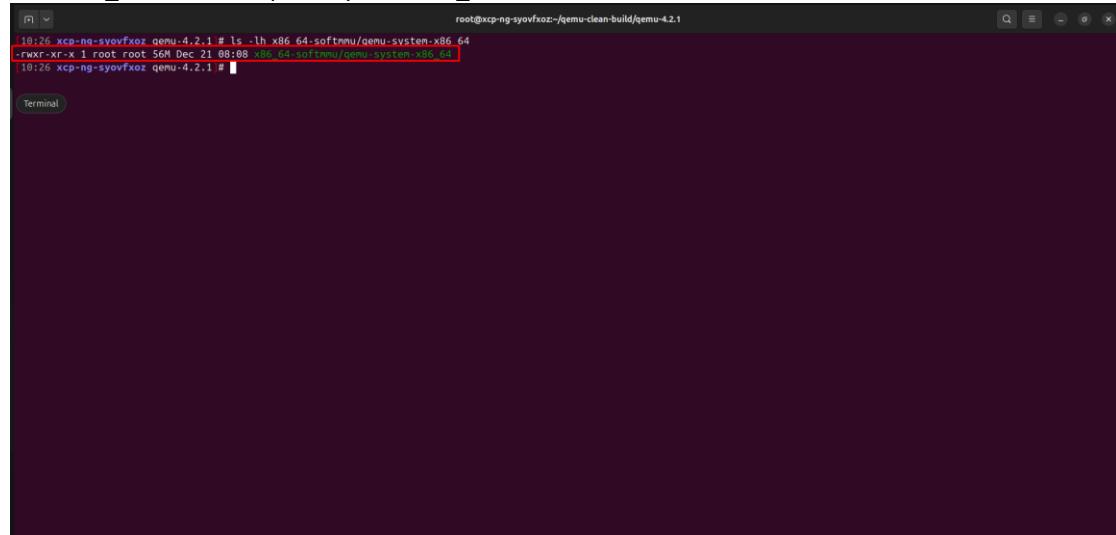
```
echo "Build exit code: $?"
# Should show: Build exit code: 0
```

If build failed:

1. Check build.log for errors: grep -i error build.log
2. Most common issue: Missing dependencies
3. Fix and run: make clean && make -j\$(nproc)

### Step 3: Verify Binary Exists

```
ls -lh x86_64-softmmu/qemu-system-x86_64
```



```
root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
10:26 xcp-ng-syovfxoz qemu-4.2.1 # ls -lh x86_64-softmmu/qemu-system-x86_64
-rwxr-xr-x 1 root root 56M Dec 21 08:08 x86_64-softmmu/qemu-system-x86_64
10:26 xcp-ng-syovfxoz qemu-4.2.1 #
```

Expected output:

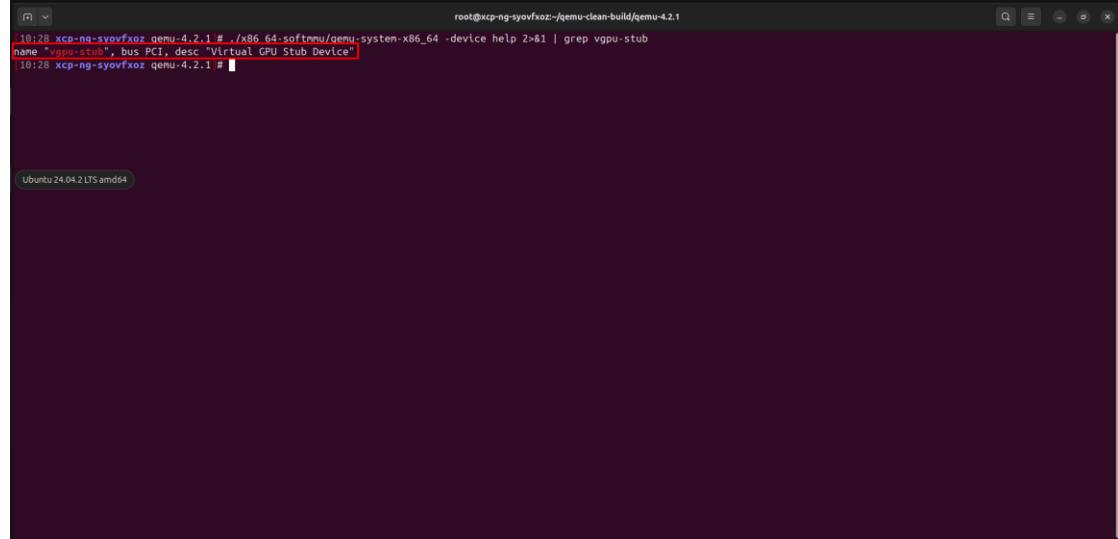
```
-rwxr-xr-x 1 root root 56M Dec 21 08:08 x86_64-softmmu/qemu-system-x86_64
```

File should be ~50-60MB and executable (rwxr-xr-x)

## SECTION 7: VERIFY THE DEVICE

### Step 1: Check Device is Available

```
./x86_64-softmmu/qemu-system-x86_64 -device help 2>&1 | grep vgpu-stub
```



```
root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
[10:28 xcp-ng-syovfxoz qemu-4.2.1] # ./x86_64-softmmu/qemu-system-x86_64 -device help 2>&1 | grep vgpu-stub
name "vgpu-stub", bus PCI, desc "Virtual GPU Stub Device"
[10:28 xcp-ng-syovfxoz qemu-4.2.1] #
```

Expected output:

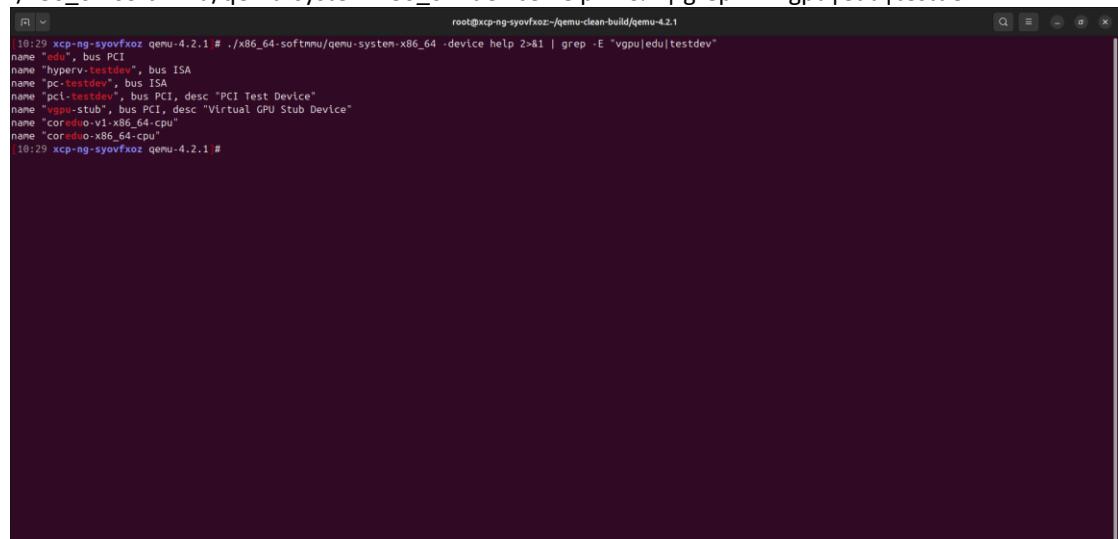
```
name "vgpu-stub", bus PCI, desc "Virtual GPU Stub Device"
```

If you don't see this, the device wasn't compiled in. Check:

```
grep vgpu-stub hw/misc/Makefile.objs
```

### Step 2: Compare with Other Test Devices

```
./x86_64-softmmu/qemu-system-x86_64 -device help 2>&1 | grep -E "vgpu|edu|testdev"
```



```
root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
[10:29 xcp-ng-syovfxoz qemu-4.2.1] # ./x86_64-softmmu/qemu-system-x86_64 -device help 2>&1 | grep -E "vgpu|edu|testdev"
name "edu", bus PCI
name "hyperv-testdev", bus ISA
name "pc-testdev", bus ISA
name "pci-testdev", bus PCI, desc "PCI Test Device"
name "vgpu-stub", bus PCI, desc "Virtual GPU Stub Device"
name "core2do-v1-x86_64-cpu"
name "core2do-x86_64-cpu"
[10:29 xcp-ng-syovfxoz qemu-4.2.1] #
```

Expected output (showing vgpu-stub along with similar devices):

```
name "edu", bus PCI
name "hyperv-testdev", bus ISA
name "pc-testdev", bus ISA
name "pci-testdev", bus PCI, desc "PCI Test Device"
name "vgpu-stub", bus PCI, desc "Virtual GPU Stub Device"
```

### Step 3: Test Instantiation

Test that the device can actually be created:

```

./x86_64-softmmu/qemu-system-x86_64 \
-M q35 \
-device vgpu-stub \
-display none \
-serial stdio \
-S &

sleep 2

ps aux | grep qemu-system-x86_64 | grep vgpu-stub

```

You should see a running QEMU process. Kill it:

```
pkill -f "qemu-system-x86_64.*vgpu-stub"
```

If QEMU starts without errors, the device is working correctly!

## **SECTION 8: DEPLOYMENT TO XCP-NG**

**CRITICAL:** XCP-ng uses /usr/lib64/xen/bin (note the lib64, not lib)

### Step 1: Find XCP-ng QEMU Binary

```
rpm -ql qemu | grep bin/qemu-system
```

Output should show:

```
/usr/lib64/xen/bin/qemu-system-i386
```

### Step 2: Verify Current QEMU Version

```
/usr/lib64/xen/bin/qemu-system-i386 --version
```

```

root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
[10:30 xcp-ng-syovfxoz qemu-4.2.1]# /usr/lib64/xen/bin/qemu-system-i386 --version
QEMU emulator version 4.2.1
Copyright (c) 2003-2019 Fabrice Bellard and the QEMU Project developers
[10:30 xcp-ng-syovfxoz qemu-4.2.1]#

```

Should show:

QEMU emulator version 4.2.1

### Step 3: Backup Original QEMU

**IMPORTANT:** Always backup before replacing!

```
cp /usr/lib64/xen/bin/qemu-system-i386 \
/usr/lib64/xen/bin/qemu-system-i386.backup-$(date +%Y%m%d-%H%M)
```

Verify backup:

```
ls -lh /usr/lib64/xen/bin/qemu-system-i386.backup-*
```

#### Step 4: Deploy New QEMU Binary

Copy our x86\_64 binary to replace qemu-system-i386:  
(The x86\_64-softmmu target can run both 32-bit and 64-bit VMs)

```
cd ~/qemu-clean-build/qemu-4.2.1
```

```
cp x86_64-softmmu/qemu-system-x86_64 /usr/lib64/xen/bin/qemu-system-i386
```

When prompted "overwrite?", type: yes

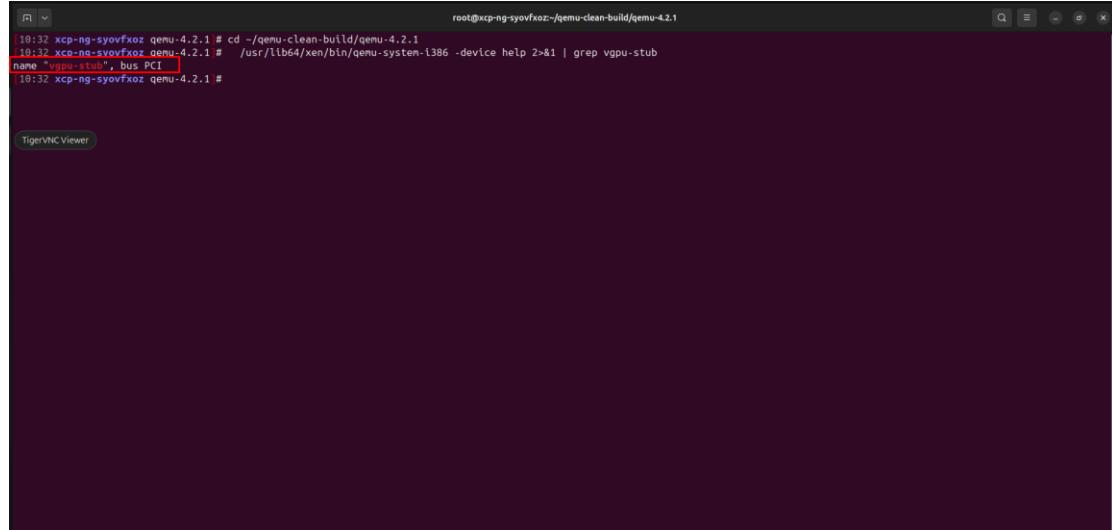
#### Step 5: Verify Deployment

Check version still matches:

```
/usr/lib64/xen/bin/qemu-system-i386 --version
```

Check device is available:

```
/usr/lib64/xen/bin/qemu-system-i386 -device help 2>&1 | grep vgpu-stub
```



```
root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1
10:32 xcp-ng-syovfxoz qemu-4.2.1 # cd ~/qemu-clean-build/qemu-4.2.1
10:32 xcp-ng-syovfxoz qemu-4.2.1 # /usr/lib64/xen/bin/qemu-system-i386 -device help 2>&1 | grep vgpu-stub
name "vgpu-stub", bus PCI
10:32 xcp-ng-syovfxoz qemu-4.2.1 #
```

Expected output:

name "vgpu-stub", bus PCI, desc "Virtual GPU Stub Device"

Check file size and permissions:

```
ls -lh /usr/lib64/xen/bin/qemu-system-i386
```

Expected:

-rwxr-xr-x 1 root root 56M Dec 21 08:21 /usr/lib64/xen/bin/qemu-system-i386

SUCCESS! QEMU with vgpu-stub is now deployed to XCP-ng!

## SECTION 9: TESTING WITH A VM

#### Method 1: Create Test VM Configuration (Simple)

Create a minimal test config:

```
cat > /tmp/test-vgpu.cfg << 'EOF'
name = "test-vgpu-stub"
type = "hvm"
memory = 1024
vcpus = 1
```

```
builder = "hvm"
boot = "c"

# Add the vgpu-stub device
device_model_args = ["-device", "vgpu-stub"]
EOF
```

Create the VM:

```
xl create /tmp/test-vgpu.cfg
```

Check if running:

```
xl list | grep test-vgpu
```

Destroy test VM:

```
xl destroy test-vgpu-stub
```

#### Method 2: Add to Existing VM

For an existing Linux VM:

1. Power off the VM:

```
xl shutdown <vm-name>
```

2. Find the VM's config:

```
ls -lh /etc/xen/
```

3. Edit the VM's config file or add parameters:

For xl-based VMs, add to config:

```
device_model_args = ["-device", "vgpu-stub"]
```

For XAPI/XCP-ng Center VMs, use xe command:

```
xe vm-param-set uuid=<vm-uuid> \
other-config:device-model-args="-device vgpu-stub"
```

4. Start the VM:

```
xl create <config-file>
```

# or

```
xe vm-start uuid=<vm-uuid>
```

## **SECTION 10: VERIFY IN GUEST VM**

#### Step 1: Access the VM

Via xl console:

```
xl console <vm-name>
```

Or via SSH if network is configured

#### Step 2: Check PCI Devices

Inside the guest VM, run:

```
lspci
```

The screenshot shows three terminal windows. The top-left window (calvin@test2) runs 'lspci' and highlights the line '01:00.0 Processing accelerator: Red Hat, Inc. Device 1111 (rev 01)'. The top-right window (david@ubuntu-root) runs 'ls -l /sys/bus/pci/devices/0000:01:00.0' and shows a 'No such file or directory' error. The bottom window (root@xcp-ng-syovfxoz) runs 'ls -l /sys/bus/pci/devices/0000:01:00.0' and lists the device node.

```

calvin@test2: ~ $ lspci
00:00.0 Host bridge: Intel Corporation 440FX - 82441FX PMC [Natoma] (rev 02)
00:01.0 ISA bridge: Intel Corporation 82371SB PIIX3 ISA [Natoma/Triton II]
00:01.1 IDE interface: Intel Corporation 82371SB PIIX3 IDE [Natoma/Triton II]
00:01.2 USB controller: Intel Corporation 82371SB PIIX3 USB [Natoma/Triton II] (rev 01)
00:01.3 VGA bridge: Intel Corporation 82371AB/EB/MB [Natoma/Triton II] (rev 01)
00:02.0 VGA compatible controller: Virtuous Logic GD 546
00:02.1 SATA controller: Virtuous Logic GD 546
00:05.0 Processing accelerators: Red Hat, Inc. Device 1111 (rev 01)
calvin@test2: ~ $ ls -l /sys/bus/pci/devices/0000:01:00.0
ls: cannot access '/sys/bus/pci/devices/0000:01:00.0': No such file or directory
calvin@test2: ~ $

root@xcp-ng-syovfxoz:~/qemu-clean-build/qemu-4.2.1 # ls -l /sys/bus/pci/devices/0000:01:00.0
01:00.0 Processing accelerator: Red Hat, Inc. Device 1111 (rev 01)

```

Expected output (look for this line):

01:00.0 Processing accelerator: Red Hat, Inc. Device 1111

You can also search specifically:

`lspci | grep -i "red hat.*1111"`

#### Step 3: Get Detailed Device Info

`lspci -v -s 01:00.0`

Expected output:

```

01:00.0 Processing accelerator: Red Hat, Inc. Device 1111
    Subsystem: Red Hat, Inc Device 1100
    Flags: fast devsel
    Memory at <address> (32-bit, non-prefetchable) [size=4K]

```

#### Step 4: Check Device Node

`ls -l /sys/bus/pci/devices/0000:01:00.0`

This shows the sysfs entry for the device.

SUCCESS! The vGPU stub device is visible in the guest VM!

## SECTION 11: TROUBLESHOOTING

### Problem: Configure fails with "xen support requested but not found"

Solution: Add `--disable-xen` to configure command (already in our guide)

### Problem: Build fails with "Python 3 required"

Solution: Add `--python=/usr/bin/python2` (already in our guide)

### Problem: Build fails with old GCC errors

Solution:

```

scl enable devtoolset-11 bash
gcc --version # Verify it shows 11.x

```

Then rebuild:

```

make clean
./configure [same options]
make -j$(nproc)

```

Problem: "common-obj-y += vgpu-stub.o" not found in Makefile.objs

Check:

```
grep -n vgpu hw/misc/Makefile.objs
```

If missing, add it manually:

```
nano hw/misc/Makefile.objs
```

Add after the edu.o line:

```
common-obj-y += vgpu-stub.o
```

Problem: Device doesn't appear in lspci in guest

Verify in sequence:

1. Check host QEMU has device:

```
/usr/lib64/xen/bin/qemu-system-i386 -device help 2>&1 | grep vgpu
```

2. Check VM is using device\_model\_args:

```
xl list -l <vm-name> | grep device_model_args
```

3. Check QEMU process has the device:

```
ps aux | grep qemu | grep <vm-name>
```

4. Check QEMU logs:

```
tail -f /var/log/xen/qemu-dm-<domain-id>.log
```

Problem: VM fails to start after deployment

Symptoms: VM won't boot, xen errors

Debugging:

1. Check Xen logs:

```
tail -100 /var/log/xen/xl-<vm>.log
```

2. Test with original QEMU:

```
cp /usr/lib64/xen/bin/qemu-system-i386.backup-* \
/usr/lib64/xen/bin/qemu-system-i386
```

3. Check QEMU permissions:

```
ls -l /usr/lib64/xen/bin/qemu-system-i386
# Should be: -rwxr-xr-x
```

4. Try without vgpu-stub first:

Remove device\_model\_args from VM config  
If VM starts, QEMU is OK but device has issues

Problem: Wrong QEMU binary path

If you have /usr/lib/xen instead of /usr/lib64/xen:

Replace all instances of /usr/lib64/xen with /usr/lib/xen in deployment steps

To find the correct path:

```
rpm -ql qemu | grep bin/qemu-system
```

## SECTION 12: ROLLBACK PROCEDURE

If you need to restore the original QEMU:

### Step 1: List Available Backups

```
ls -lh /usr/lib64/xen/bin/qemu-system-i386.backup-*
```

Step 2: Restore Backup

```
# Replace TIMESTAMP with actual backup timestamp  
cp /usr/lib64/xen/bin/qemu-system-i386.backup-TIMESTAMP /usr/lib64/xen/bin/qemu-system-i386
```

Step 3: Verify

```
/usr/lib64/xen/bin/qemu-system-i386 --version  
/usr/lib64/xen/bin/qemu-system-i386 -device help 2>&1 | grep vgpu
```

The vgpu-stub should no longer appear.